Bhanuprasad Kokkula

New York, NY | (848) 391-1376 | bhanu.kokkula457@gmail.com
linkedin.com/in/bhanu-kokkula | github.com/Bhanukokkula | bhanukokkula.github.io

EDUCATION & CERTIFICATIONS

Rutgers, The State University of New Jersey — M.S. in Data Science Aug 2024 – May 2026

University of Pennsylvania (Wharton School) — Professional Certificate: AI for Business Jan 2026 – Jun 2026

Certifications: Bloomberg Market Concepts · Microsoft Certified: Azure Data Scientist Associate

Selected Coursework: Machine Learning Summer Program – IIT Kharagpur

EXPERIENCE

Rutgers, The State University of New Jersey — Graduate GenAI Engineer Oct 2024 – May 2026

- Owned an end-to-end retrieval pipeline that ingests, cleans, and indexes **100K+** documents into a searchable knowledge base, tuning retrieval to sub-**200**ms at production query loads.
- Built an automated evaluation harness benchmarking **15+** model configurations on accuracy, factual grounding, and latency, cutting evaluation cycles from **2** weeks to **3** hours and turning model selection into a data-driven decision rather than a judgment call.
- Ran controlled experiments across **5K+** test cases to isolate retrieval vs. tuning effects, identified the configuration lifting answer accuracy **31%** over baseline, and shipped it as the team default.

Hexaware Technologies — Data Scientist, Financial Services Jun 2023 – Aug 2024

- Owned the feature layer behind production fraud-detection models, building and validating **12** feature pipelines over **500K+** daily transactions; those features drove model decisions that cut fraud losses **8%.**
- Built a credit-risk feature pipeline over **5M+** borrower records with automated validation gates and retraining triggers, turning raw data into model-ready inputs that held up under live production drift.
- Designed statistical monitoring across **50+** fields to catch distribution shifts and silent data corruption before they reached downstream models, surfacing data-quality issues the business had previously absorbed blind.
- Automated **14** manual reporting workflows into validated pipelines, cutting turnaround from **3** days to **4** hours.

Chinmaya Technologies — Data Science Intern Jan 2023 – May 2023

- Engineered **40+** behavioral features from **1M+** transaction and clickstream records and built SKU-level demand forecasts that cut stockouts **12%.**

Echortech — Data Engineering Intern Sep 2022 – Dec 2022

- Built and deployed ETL pipelines moving millions of records daily from **6** source systems into BigQuery, with automated data-quality checks that caught schema drift and null-rate spikes previously going undetected.

RESEARCH & PUBLICATION

Rutgers, The State University of New Jersey — Research Co-author, Computational Linguistics (ongoing) Jun 2025

- Engineered a quantitative text-mining framework executing semantic claim-matching across **194K+** abstracts and **69,580** full-text papers to isolate a **46.3%** baseline rate of narrative overreach, applying journal-clustered logistic regressions and pre-registered behavioral trials (N = **1,105**) to evaluate statistical claim inflation.

SELECTED PROJECTS

GLASSBOX — Cross-Sectional Equity Factor Backtester Python

- Built an event-driven backtester from scratch for cross-sectional equity factor research, engineered to resist four classic failure modes (lookahead bias, survivorship, overfitting, multiple testing) that inflate paper returns.
- Validated a momentum factor on point-in-time data, reporting a net Sharpe of **0.35** and a Deflated Sharpe Ratio of **0.96** to discount for multiple testing, backed by **89** passing tests for reproducibility.

MARS — Multi-Source Alternative-Data Research System Python, LangGraph

- Built a research system ingesting and synthesizing SEC **10-K**/**10-**Q filings, macroeconomic indicators, and real-time news across **50K+** financial documents, extracting entities, scoring sentiment, and flagging cross-source anomalies for analyst review.
- Designed a conflict-resolution protocol reconciling contradictory signals across earnings, news, macro, and filings, reaching **68%** directional accuracy on forward-looking risk flags and cutting per-company research from **6** hours to **15** minutes.

Causal Inference & Uplift Modeling Python

- Built a causal-inference framework on **2M+** observational records, using propensity-score matching to construct a balanced synthetic control and cut covariate imbalance below **0.1** across **25+** features, recovering experiment-grade estimates from non-experimental data.
- Estimated treatment effects with difference-in-differences and two-way fixed effects (**4.7%** effect, p < **0.01**), then applied uplift modeling to isolate the top-**30%** responsive segment for **2.3x** higher return than a blanket approach.

SKILLS

Programming: Python, R, SQL, Spark, Bash

Statistics & Causal Inference: Hypothesis & A/B Testing, Propensity-Score Matching, Difference-in-Differences, Uplift Modeling, Time-Series, Bayesian Inference

Machine Learning & AI: Scikit-Learn, XGBoost, LightGBM, PyTorch, MLflow, LangGraph, Hugging Face, RAG

Data & Cloud: PostgreSQL, Snowflake, BigQuery, Airflow, Docker, Git, FastAPI; AWS (SageMaker, S3, Lambda), GCP, Azure